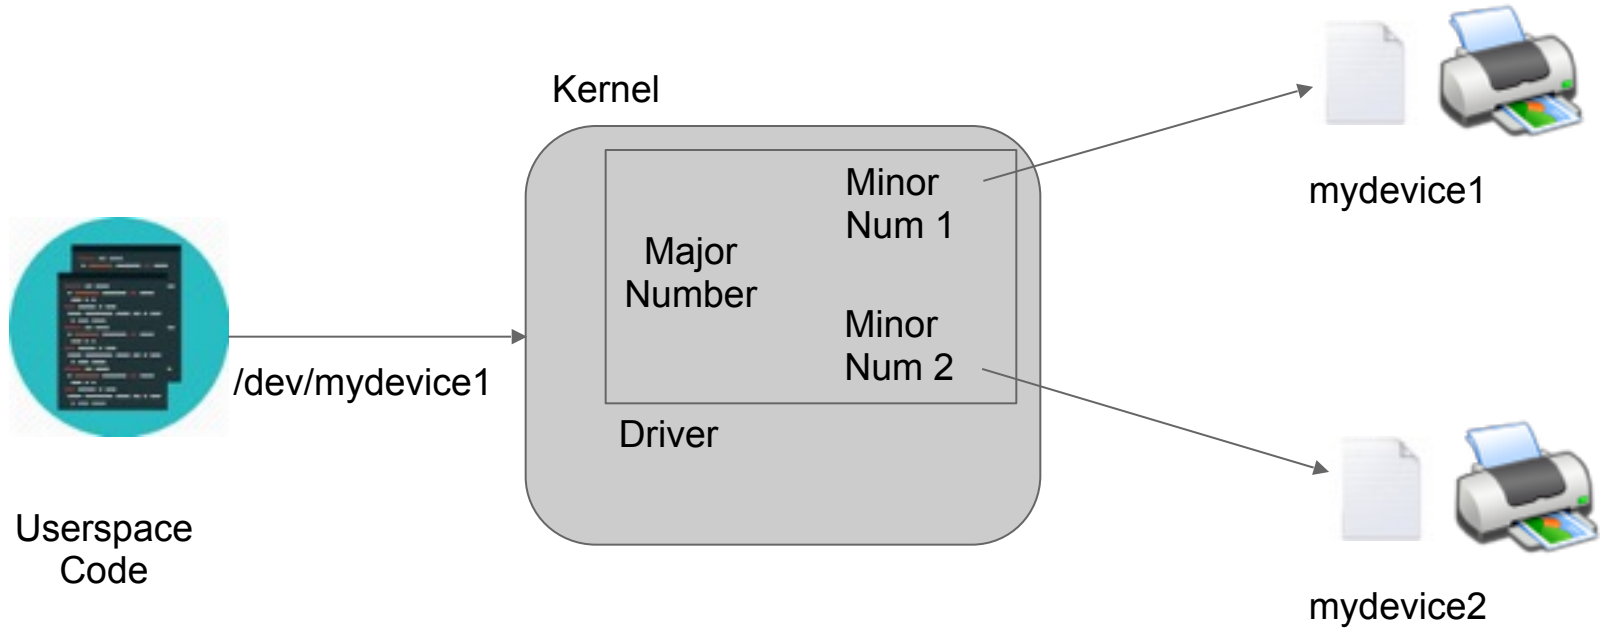


CS350 Lab5

Major and Minor Number



Step1: Write kernel_module.c

- Steps given in kernel module slides.
- Kernel module should have a file name
 - E.g. “mydevice” as given in slides
- It should define the allowed file operations on this file
 - “my_fops” in the slides contains the function pointers to the allowed file operation functions
- Driver should also request a minor number for the device
 - using MISC_DYNAMIC_MINOR
 - operating system dynamically assigns minor number to this file
- To register this device with kernel, you must call “misc_register()” function.
 - Ideal place to call “ misc_register()” function is in init_module() function as it is the first function that is called when you insert the module in the kernel.
- To unregister the device, you should call “misc_deregister()”.
 - Ideal place to call it is in cleanup_module()

Step 2: Write user space program

- We can use regular file operations on device files:
 - Open - Called each time the device is opened from user space.
 - Read - Called when a process has already opened the file and tries to read from it.
 - Use “copy_to_user” to copy data to user space from kernel
 - Write - Called when a process tries to write into the device file.
 - Use “copy_from_user” to copy data to kernel from user space
 - Close - Called when the device is closed in user space.

How do file ops work on character devices

- A file operation on a device file will be handled by the kernel module associated with the device.
- Call “open” system call to open “mydevice” file
- Call “read” system call to read from the “mydevice” file

```
fd = open("/dev/mydevice", O_RDWR);
```

- opens /dev/mydevice device for read and write operation.
- OS will call my_open() file operation handler in the kernel module which is associated with the device.
- misc_register(&my_misc_device) instruction in my_module_init() registers the module. It creates an entry in the “/dev” directory for “mydevice” file and informs the operating system what file-operations handler functions are available for this device.

Memory allocation/deallocation in Kernel

Memory Allocation:

kmalloc(): Allocates physically contiguous memory

`void * kmalloc(size_t size, int flags)`

kzalloc(): Allocates memory and sets it to zero

vmalloc(): Allocates memory that is virtually contiguous and not necessarily physically contiguous.

`void * vmalloc(unsigned long size)`

Memory Deallocation: kfree()

Moving data in and out of the Kernel

- **copy_to_user()**

- unsigned long copy_to_user (void __user * *dst*, const void * *src*, unsigned long *n*);
- Copies data **from kernel space to user space**
- Returns number of bytes that could not be copied. On success, this will be zero.
- Checks that *dst* is writable by calling access_ok on *dst* with a type of VERIFY_WRITE. If it returns non-zero, copy_to_user proceeds to copy

- **copy_from_user()**

- unsigned long copy_from_user (void * *dst*, const void __user * *src*, unsigned long *n*);
- Copies data **from user space to kernel**
- Returns number of bytes that could not be copied. On success, this will be zero.

Question: Why shouldn't you use **memcpy** or **call by reference** to access userspace data?