


## Major and Minor Number



# Step1: Write kernel\_module.c

- Steps given in kernel module slides.
- Kernel module should have a file name
  - E.g. “mydevice” as given in slides
- It should define the allowed file operations on this file
  - “my\_fops” in the slides contains the function pointers to the allowed file operation functions
- Driver should also request a minor number for the device
  - using MISC\_DYNAMIC\_MINOR
  - operating system dynamically assigns minor number to this file
- To register this device with kernel, you must call “misc\_register()” function.
  - Ideal place to call “ misc\_register()” function is in init\_module() function as it is the first function that is called when you insert the module in the kernel.
- To unregister the device, you should call “misc\_deregister()”.
  - Ideal place to call it is in cleanup\_module()

# Step 2: Write user space program

- We can use regular file operations on device files:
  - Open - Called each time the device is opened from user space.
  - Read - Called when a process has already opened the file and tries to read from it.
    - Use “copy\_to\_user” to copy data to user space from kernel
  - Write - Called when a process tries to write into the device file.
    - Use “copy\_from\_user” to copy data to kernel from user space
  - Close - Called when the device is closed in user space.

# How do file ops work on character devices

- A file operation on a device file will be handled by the kernel module associated with the device.
- Call “open” system call to open “mydevice” file
- Call “read” system call to read from the “mydevice” file

```
fd = open("/dev/mydevice", O_RDWR);
```

- opens /dev/mydevice device for read and write operation.
- OS will call my\_open() file operation handler in the kernel module which is associated with the device.
- misc\_register(&my\_misc\_device) instruction in my\_module\_init() registers the module. It creates an entry in the “/dev” directory for “mydevice” file and informs the operating system what file-operations handler functions are available for this device.

# Memory allocation/deallocation in Kernel

## Memory Allocation:

`kmalloc()`: Allocates physically contiguous memory

`void * kmalloc(size_t size, int flags)`

`kzalloc()`: Allocates memory and sets it to zero

`vmalloc()`: Allocates memory that is virtually contiguous and not necessarily physically contiguous.

`void * vmalloc(unsigned long size)`

## Memory Deallocation: `kfree()`

# Moving data in and out of the Kernel

- **copy\_to\_user()**

- unsigned long copy\_to\_user (void \_\_user \* *dst*, const void \* *src*, unsigned long *n*);
- Copies data **from kernel space to user space**
- Returns number of bytes that could not be copied. On success, this will be zero.
- Checks that *dst* is writable by calling access\_ok on *dst* with a type of VERIFY\_WRITE. If it returns non-zero, copy\_to\_user proceeds to copy

- **copy\_from\_user()**

- unsigned long copy\_from\_user (void \* *dst*, const void \_\_user \* *src*, unsigned long *n*);
- Copies data **from user space to kernel**
- Returns number of bytes that could not be copied. On success, this will be zero.

**Question:** Why shouldn't you use **memcpy** or **call by reference** to access userspace data?