

Using semaphores and waitqs in kernel

# Semaphores

## 1. Declaration of Semaphore Variables

```
static DEFINE_SEMAPHORE(semVar); // Declare and initialize a semaphore to 1
```

## 2. Initialization of counting semaphores

```
sema_init(&semVar, 3); // semVar initialized to 3
```

## 3. Use

```
down_interruptible(&semVar); // DON'T USE down(&semVar1)...why?  
up(&semVar);
```

Other versions of down operations can be found here:

<http://lxr.free-electrons.com/source/include/linux/semaphore.h>

<http://lxr.free-electrons.com/source/kernel/locking/semaphore.c>

- `down_killable`
- `down_trylock`
- `down_timeout`

# Mutex

- initialize the mutex as unlocked

```
static DEFINE_MUTEX(mutVar);
```

- re-initialize the mutex as unlocked

```
mutex_init(&mutVar)
```

- Usage

```
mutex_lock_interruptible(&mutVar);  
//Critical Section  
mutex_unlock(&mutVar);
```

- More versions of mutex\_lock can be found here:

- <http://lxr.free-electrons.com/source/include/linux/mutex.h>
- <http://lxr.free-electrons.com/source/kernel/locking/mutex.c>
- mutex\_lock\_killable
- mutex\_lock\_trylock

# Sample Code

```
static DEFINE_SEMAPHORE(sem1);
static DEFINE_SEMAPHORE(sem2);
static DEFINE_MUTEX(mut);

//inside init_module() function
sema_init(&sem1, N);
sema_init(&sem2, 0);
mutex_init(&mut); //redundant because DEFINE_MUTEX also initializes mut

...
down_interruptible(&sem1);
mutex_lock_interruptible(&mut);
Critical Section
mutex_unlock(&mut);
up(&sem2);
...
```

**IMPORTANT:** Don't forget to check for error return from \*\_interruptible() functions.  
**Why?**

# wait queues (waitq)

- See: <http://lxr.free-electrons.com/source/include/linux/wait.h>
- A waitqueue is a queue of processes that are waiting for a specific event.
- Declaration  
`static wait_queue_head_t wq;`
- Initialization  
`init_waitqueue_head(&wq);`
- OR Declare and Initialize in one shot  
`static DECLARE_WAIT_QUEUE_HEAD(wq);`
- Making the current process wait for a condition to be true  
`wait_event_interruptible(wq, condition);`
- Preferred over `wait_event(...)`
- Waking up a waiting process upon an event  
`wake_up(&wq)`
  - wake up all processes waiting on this wait queue. Woken processes check the condition. If condition is false, they go back to sleep.
- `wake_up_interruptible(&wq)`
  - wakes up only the processes that are in interruptible waits. Any process that in non-interruptible waits will continue to sleep.

# Sample code: using mutex and waitqs together

```
static DECLARE_WAIT_QUEUE_HEAD(wq);  
static DEFINE_MUTEX(mut);
```

```
mutex_lock_interruptible(&mut); //enter critical section
```

```
...
```

```
//wait for an event somewhere in the critical section
```

```
while(!condition) {
```

```
    //release mutex as you can't sleep holding it
```

```
    mutex_unlock(&mut);
```

```
    //sleep on condition
```

```
    wait_event_interruptible(wq,condition);
```

```
    //reacquire mutex before re-entering critical section
```

```
    mutex_lock_interruptible(&mut);
```

```
    //now go back check condition again.
```

```
    //Why?
```

```
    //because return from wait and re-locking are not atomic
```

```
    //So condition may have changed again
```

```
}
```

```
...
```

```
mutex_unlock(&mut)//exit critical section
```

- And someone has to wake up the sleepers

```
wake_up(&wq);
```

OR

```
wake_up_all(&wq);
```

# Useful Header Files

```
#include <linux/module.h>
#include <linux/miscdevice.h>
#include <linux/fs.h>
#include <asm/uaccess.h>
#include <linux/slab.h>

#include <linux/wait.h>
#include <linux/semaphore.h>
#include <linux/mutex.h>
```