

# **CS350 Lab7**

# Adding a “new” System call to Kernel

You can do it in two ways:

1. Easy and smart approach:

Add the new system call by modifying an **existing** kernel file

2. Generic approach:

Add the new system call by creating a **new** file in kernel

In following slides, generic approach is used to add a simple system call

# Generic Approach

## Step 1:

Create an entry for the system call in the kernel's syscall\_table

- This file is located at ~/linux-4.2.5/source/arch/x86/entry/syscall/**syscall\_64.tbl**
- Add a new entry for this new system call under “**64-bit system call numbers and entry vectors**” section. The number assigned for this system call must be unique.

In syscall\_64.tbl for linux-4.2.5 version, you will notice 0 to 322 numbers have already been used by other system calls, so add an entry with a new number.

e.g.

323	common	myfoo	sys_myfoo
-----	--------	-------	-----------

# Adding a new entry in syscall\_64.tbl

Adding a new entry at the end of the list of system calls under “**64-bit system call numbers and entry vectors**” heading e.g.

```
323 common myfoo sys_myfoo
```

```

318      common  getrandom      sys_getrandom
319      common  memfd_create   sys_memfd_create
320      common  kexec_file_load sys_kexec_file_load
321      common  bpf            sys_bpf
322      64      execveat       stub_execveat
323      common  myfoo          sys_myfoo
#
# x32-specific system call numbers start at 512 to avoid cache impact
# for native 64-bit operation.
#
512      x32     rt_sigaction   compat_sys_rt_sigaction
513      x32     rt_sigreturn   stub_x32_rt_sigreturn
514      x32     ioctl          compat_sys_ioctl
515      x32     readv           compat_sys_readv
516      x32     writev          compat_sys_writev

```

# Generic Approach

## Step 2.a:

Write the system call code as a kernel function i.e. implementation of this new system call

- To do that (In generic approach), create a hello.c with following content and place this hello.c file in ~/linux-4.2.5/kernel/ folder:

```
#include <linux/linkage.h>
#include <linux/export.h>
#include <linux/time.h>
#include <asm/uaccess.h>
#include <linux/printk.h>
#include <linux/slab.h>

asmlinkage int sys_myfoo(void){
    printk(KERN_ALERT "Hello World!\n");
    return 0;
}

EXPORT_SYMBOL(sys_myfoo);
```

# Generic Approach

## Step 2.b:

Update the “Makefile” in ~/linux-4.2.5/kernel/ folder so “hello.c” gets compiled when you give command to compile the whole kernel.

- To update this Makefile: just add “hello.o” at the end of the list of object file names.

```
#  
# Makefile for the linux kernel.  
#  
  
obj-y      = fork.o exec_domain.o panic.o \  
             cpu.o exit.o softirq.o resource.o \  
             sysctl.o sysctl_binary.o capability.o ptrace.o user.o \  
             signal.o sys.o kmod.o workqueue.o pid.o task_work.o \  
             extable.o params.o \  
             kthread.o sys_ni.o nsproxy.o \  
             notifier.o ksysfs.o cred.o reboot.o \  
             async.o range.o smpboot.o hello.o  
  
obj-$(CONFIG_MULTIUSER) += groups.o  
  
ifdef CONFIG_FUNCTION_TRACER  
# Do not trace debug files and internal ftrace files
```

# Adding a “new” System call to Kernel

## Step 3.a:

- Compile your kernel i.e. follow the instructions given in this link:

[http://www.cs.binghamton.edu/~kartik/cs350/lab\\_slides/kernel\\_compilation.html](http://www.cs.binghamton.edu/~kartik/cs350/lab_slides/kernel_compilation.html)

If you have already successfully compiled it before adding this new system call, then follow the instructions from step 7, else follow the instructions from step 1.

- After successful compilation, you will see “hello.o” object file is created in ~/linux-4.2.5/kernel/ folder.
- After compiling and rebooting your machine with your new image successfully, you can use this new system call in user space (see next slide).

# Invoke your new handler with syscall

## Step 3.b:

- Use the syscall() library function as explained in the system\_calls.pdf slides

```
#include<stdio.h>
```

```
#include<unistd.h>
```

```
#include<linux/unistd.h>
```

```
#include<linux/time.h>
```

```
int main(){
```

```
    int y = 2;
```

```
    y = syscall(323);
```

```
    printf("syscall return value :%d\n",y);//negative value of y will indicate a failure
```

```
    return 0;
```

```
}
```

- Successful call will print “Hello World!” in /var/log/kern.log file
- Run “dmesg” command to check the content of the file