


CPU Scheduling

Section 2.4: Tanenbaum's book
Chapter 5: Silberschatz's book

Kartik Gopalan


Process Lifecycle Revisited



1. Process blocks for input
2. Scheduler picks another process
3. Scheduler picks this process
4. Input becomes available

- Ready
 - Process is ready to execute, but not yet executing
 - Its waiting in the scheduling queue for the CPU scheduler to pick it up.
- Running
 - Process is executing on the CPU
- Blocked
 - Process is waiting (sleeping) for some event to occur.
 - Once the event occurs, process will be woken up, and placed on the scheduling queue.

CPU scheduler



- Selects the next process to run on the CPU from among the processes that are ready to execute
- CPU scheduling decision may take place when any process:
 1. Switches from running to waiting state
 2. Switches from running to ready state (pre-emptive scheduling)
 3. Switches from waiting to ready
 4. Terminates

Dispatcher

- Not the same as scheduler.
- Dispatcher gives control of the CPU to the process selected by the scheduler; this involves:
 - switching context
 - switching to user mode
 - jumping to the proper location in the user program to restart that program
- *Dispatch latency*
 - Time it takes for the dispatcher to stop one process and start another running

Linux CPU Scheduling Code

- To browse, go to <http://lxr.free-electrons.com>
- To download, modify, and compile, go to
 - <http://ftp.kernel.org/>
- Scheduling code is located under
 - `kernel/sched/core.c`
- Three ways to invoke kernel code
 - System calls
 - Hardware interrupts
 - Exceptions
- For x86, all entry points into the 4.2 kernel are located in assembly code in
 - `arch/x86/entry/entry_32.S`
 - OR `arch/x86/entry/entry_64.S`

When does the scheduler execute?

A. Timer interrupt fires OR

B. Some process blocks (or gives up CPU)

- Look at the CPU scheduler code in Linux
 - `schedule()` function in `kernel/sched/core.c`
 - <http://lxr.free-electrons.com/source/kernel/sched/core.c#L2988>
- Entry point to the kernel (in `entry_*.S`) invokes `schedule()` function which then
 - Figures out the next process to schedule and
 - Swaps the prev and next process via `context_switch()` function, which in turn
 - Switches memory state using `switch_mm()` and
 - Switches register+stack state using `switch_to()`

Scheduling Criteria

- **CPU utilization** – keep the CPU as busy as possible
 - **Throughput** – Number of processes that complete their execution per time unit
 - **Turnaround time** – amount of time to execute a particular process, from submission to termination.
 - **Waiting time** – amount of time a process has been waiting in the ready queue
 - **Response time** – amount of time it takes from when a request was submitted until the first response is produced.
- ## Optimization Criteria
- Max CPU utilization
 - Max throughput
 - Min turnaround time
 - Min waiting time
 - Min response time

Different systems have different scheduling goals

All systems

Fairness - giving each process a fair share of the CPU

Policy enforcement - seeing that stated policy is carried out

Balance - keeping all parts of the system busy

Batch systems

Throughput - maximize jobs per hour

Turnaround time - minimize time between submission and termination

CPU utilization - keep the CPU busy all the time

Interactive systems

Response time - respond to requests quickly

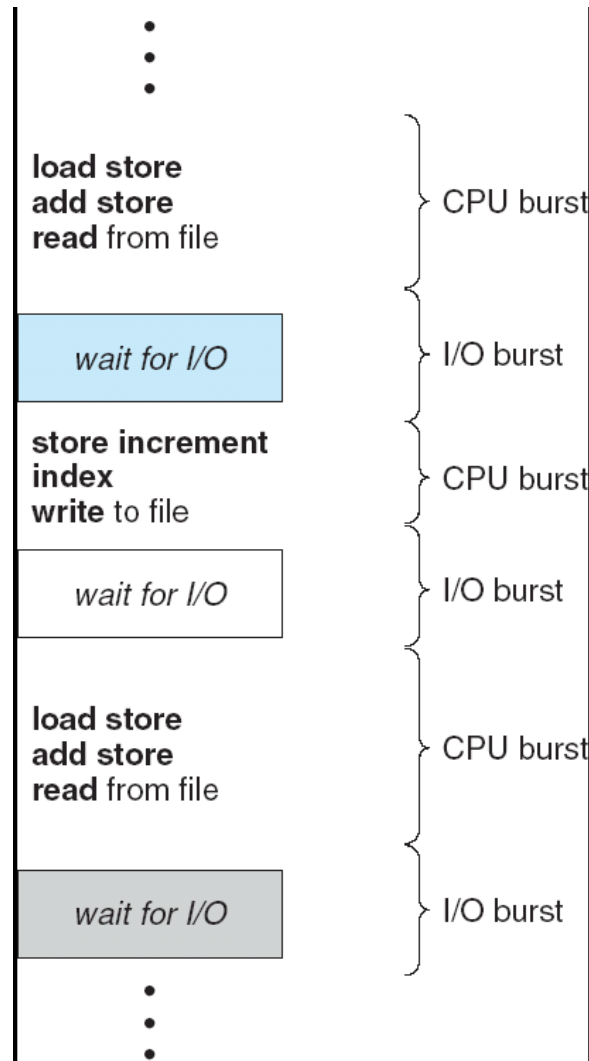
Proportionality - meet users' expectations

Real-time systems

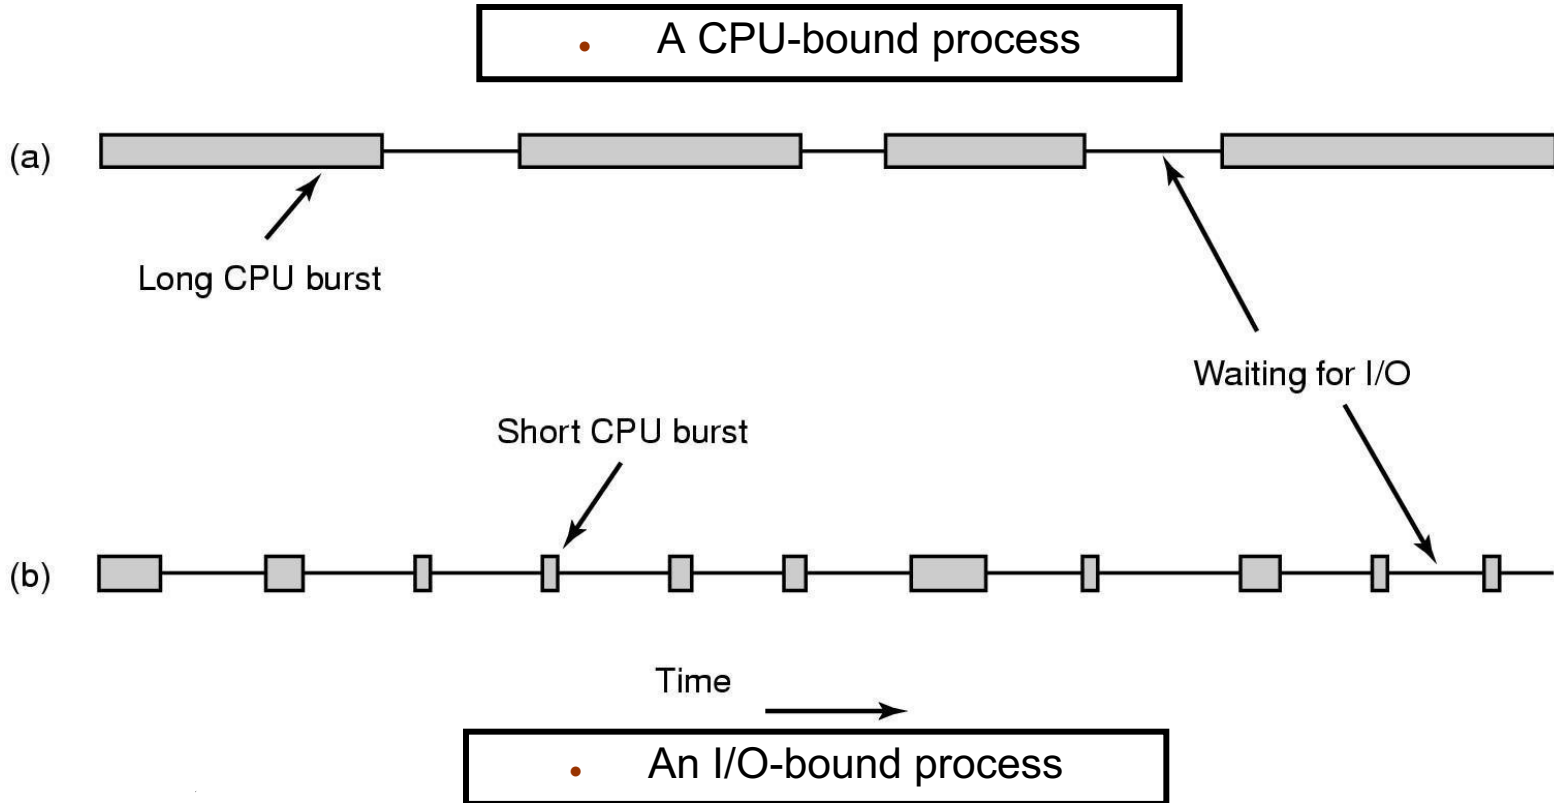
Meeting deadlines - avoid losing data

Predictability - avoid quality degradation in multimedia systems

Alternating Sequence of CPU And I/O Bursts in a Typical Program



Alternating Sequence of CPU And I/O Bursts (contd)




First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

Suppose that the processes arrive in the order: P_1 , P_2 , P_3 . The Gantt Chart for the schedule is:

- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$




FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order

$$P_2, P_3, P_1$$

- The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- *Convoy effect* : short process behind long process


Shortest-Job-First (SJF) Scheduling

- Consider the length of its next CPU burst for each process.
- Schedule the process having the shortest next CPU burst.
- Two schemes:
 - **nonpreemptive** – once CPU given to the process it cannot be preempted until completes its CPU burst
 - **preemptive** – if a new process arrives with CPU burst length less than remaining time of currently executing process, preempt the current process. This scheme is known as the Shortest Remaining Time First (SRTF)
- **Claim: SJF algorithm is optimal**
 - SJF Gives a schedule with least average waiting time compared to any possible scheduling algorithm.

Example of Non-Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

- SJF (non-preemptive)




- Average waiting time = $(0 + 6 + 3 + 7)/4 = 4$

Example of Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0.0	7
P_2	2.0	4
P_3	4.0	1
P_4	5.0	4

- SJF (preemptive)



- Average waiting time = $(9 + 1 + 0 + 2)/4 = 3$

SJF is Optimal w.r.t average wait time

- Meaning that no other algorithm can achieve a lower wait time than SJF

Proof:

- Assume that there was an algorithm X that gave better average wait time than SJF for a set of N processes.
- Since X is not the same as SJF, it means that there must be at least two processes P1 and P2 in the schedule generated by X such that
 - a. P1 executes before P2 does and
 - b. The CPU execution time of P1 is longer than that of P2 and
 - c. The average wait time of the schedule is smaller than that given by SJF
- BUT, if you swap the positions of P1 and P2 in the schedule generated by X, then the average wait time goes down!
- So keep swapping all such process pairs that satisfy conditions (a) and (b) above. Each swap will reduce the average wait time.
- Finally you will end up with a schedule generated by SJF, whose wait time cannot be reduced any further. Hence SJF is Optimal.

Exponential Averaging: Determining the Length of Next CPU Burst

- Not easy. Can only *guess* the length of next CPU burst
- Can be done by using the length of previous CPU bursts, using **exponential averaging**

t_n = actual length of the n^{th} CPU burst


τ_n = predicted value of the n^{th} CPU burst

α , $0 \leq \alpha \leq 1$

Define:

$$\tau_{n+1} = \alpha t_n + (1-\alpha) * \tau_n$$

Prediction of the Length of the Next CPU Burst



CPU burst (t_i)	6	4	6	4	13	13	13	...
"guess" (τ_i)	10	8	6	6	9	11	12	...

Examples of Exponential Averaging

$$\tau_{n+1} = \alpha t_n + (1-\alpha) * \tau_n$$

- $\alpha = 0$
 - $\tau_{n+1} = \tau_n$
 - CPU burst history does not count
- $\alpha = 1$
 - $\tau_{n+1} = \alpha t_n$
 - Only the actual last CPU burst counts
- If we expand the formula, we get:
$$\begin{aligned}\tau_{n+1} = & \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ & + (1 - \alpha)^{n+1} \tau_0\end{aligned}$$
- Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor


Round Robin (RR)

- Each process gets a fixed unit of CPU burst time (*time quantum*)
 - usually 10 to 100 milliseconds.
- After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in bursts of at most q time units at once. No process waits more than $(n-1)q$ time units.
- Performance
 - q large \Rightarrow FIFO, because processes rarely get pre-empted
 - q small \Rightarrow Smaller response times, but too much context switching overhead. q must be large compared to context switch time, otherwise overhead is too high

Example of RR with Time Quantum = 20


<u>Process</u>	<u>Burst Time</u>
P_1	53
P_2	17
P_3	68
P_4	24

- The Gantt chart is:



- Typically, higher average turnaround than SJF, but
 - Better *response time*
 - No Starvation

Time Quantum and Context Switch Time

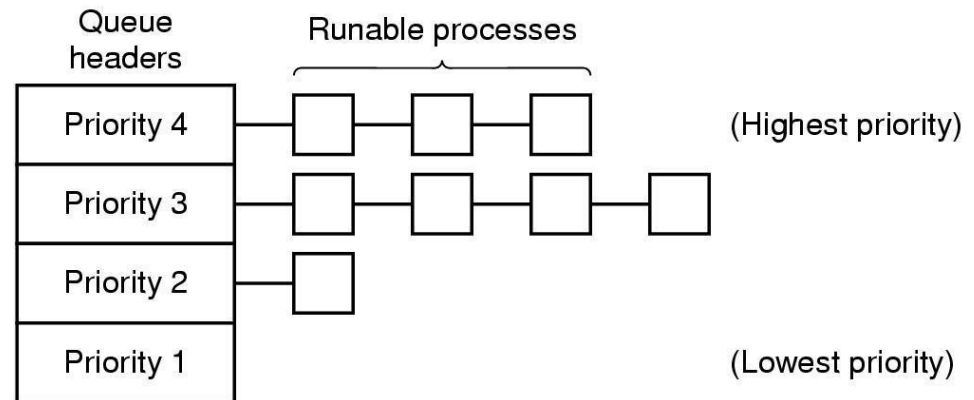


- Smaller the time quantum, more the number of context switches.
- Larger the time quantum, larger the response time.
- Scheduling algorithm needs to find a balance between context switch overhead and response time.

Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)

- Again two types
 - Preemptive
 - nonpreemptive

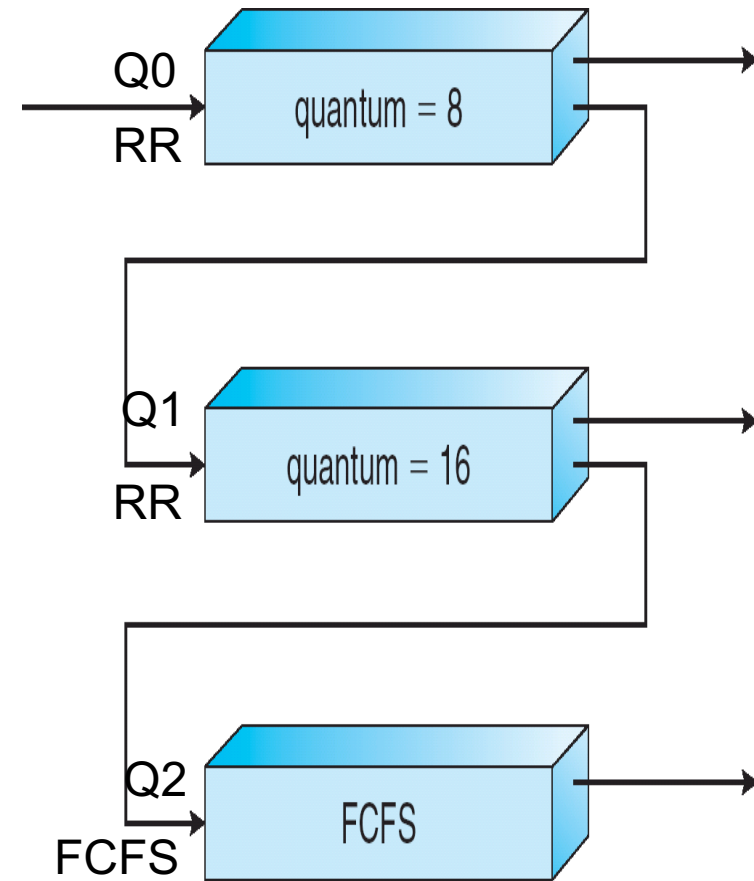


Example with four priority classes

- SJF is a priority scheduling algorithm where priority is the predicted next CPU burst time
- Problem \equiv **Starvation** \rightarrow low priority processes may never execute
- Solution \equiv **Aging** \rightarrow as time progresses increase the priority of a lower priority process that is not receiving CPU time.

Multilevel Feedback Queue (MFQ)

- Ready queue is partitioned into separate queues
- Each queue has its own scheduling algorithm
- A process can move between the various queues;
- MFQ scheduler defined by :
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade or demote a process
- Example of MFQ that gives higher priority to interactive jobs
 - A new job enters queue Q_0 which is served RR.
 - When it gains CPU, it receives 8 ms. If it doesn't finish in 8 ms, job is moved to Q_1 .
 - At Q_1 job is again served RR; receives 16 ms.
 - If it still doesn't complete, it is preempted and moved to queue Q_2 where it is served FCFS.



Real-Time Scheduling

- *When each task needs to be completed before a given deadline*
- *Hard real-time systems*
 - Required to complete a critical task before its deadline
 - For example, in a flight control system, or nuclear reactor
- *Soft real-time systems*
 - Meeting deadlines desirable, but not essential
 - For example, video or audio
- *Schedulability criteria*
 - Given m periodic events, where event i occurs within period P_i and requires C_i computation time each period
 - Then the load can be handled only if
$$\sum_i (C_i/P_i) \leq 1$$
 - The above condition is *necessary* but *NOT sufficient*.

Fair Scheduling

- Notion of “fairness” does not necessarily mean equal CPU share for all processes.
- Say you have N processes
- Each process P_i is assigned a weight w_i
- The CPU time will be divided among processes in proportion to their weights.
- Let's say some process does not use its assigned CPU time.
 - The “spare” CPU time is divided among the remaining ready processes according to the ratio of their weights.

Work-conserving versus non-work-conserving

- Work-conserving scheduler
 - CPU will not remain idle if there are processes in the ready queue
- Non-work-conserving scheduler
 - Under some conditions, scheduler may decide to “waste” CPU time even though there may be processes sitting in the ready queue
 - E.g. Sometimes real-time process cannot be started before a given time (release time).

Linux CPU Scheduling Code

To browse, go to <http://lxr.linux.no/>

To download, modify, and compile, go to
<http://ftp.kernel.org/>

Scheduling code is located under
`kernel/sched.c`

Entry points into the kernel are located at
`arch/x86/entry_32.S`
OR `arch/x86/entry_64.S`