


Memory Management Basics

From Chapter 4, Modern Operating Systems, Andrew S. Tanenbaum


Swapping (1)




- ◆ Physical memory may not be enough to accommodate needs of all processes
- ◆ Memory allocation changes as
 - processes come into memory
 - leave memory and are **swapped out** to disk
 - Re-enter memory by getting **swapped-in** from disk
- ◆ Shaded regions are unused memory

Virtual Memory

- ◆ Swapping is useful when the sum total of memory requirements of all processes is greater than DRAM available in the system.
- ◆ But sometimes, a single process might require more memory than the available DRAM in the system.
- ◆ In such cases swapping is not enough. Rather, we need to break up the memory space of a process into smaller equal-sized pieces (called pages).
- ◆ Operating system then decides which pages stay in memory and which get moved to disk.
- ◆ Virtual memory: means that each process gets an illusion that it has more memory than the physical DRAM in the system.



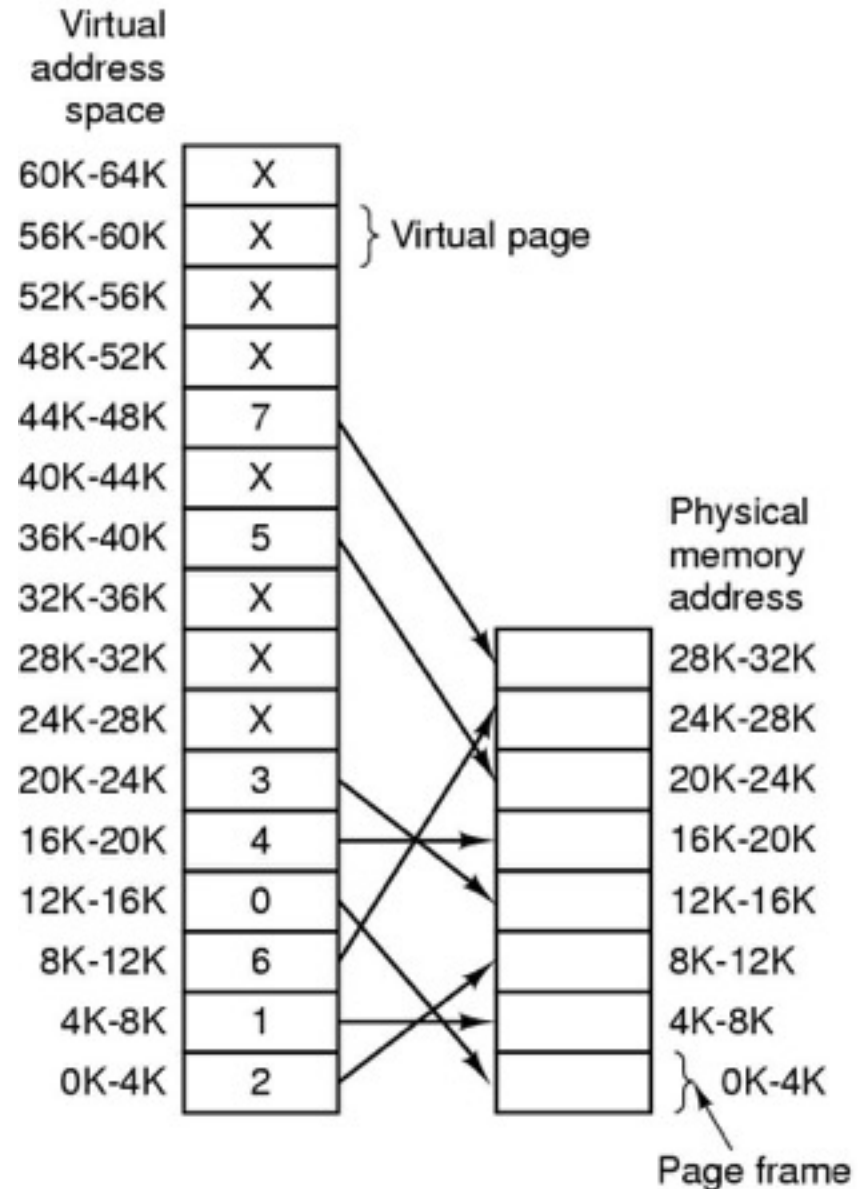
Virtual Memory and MMU



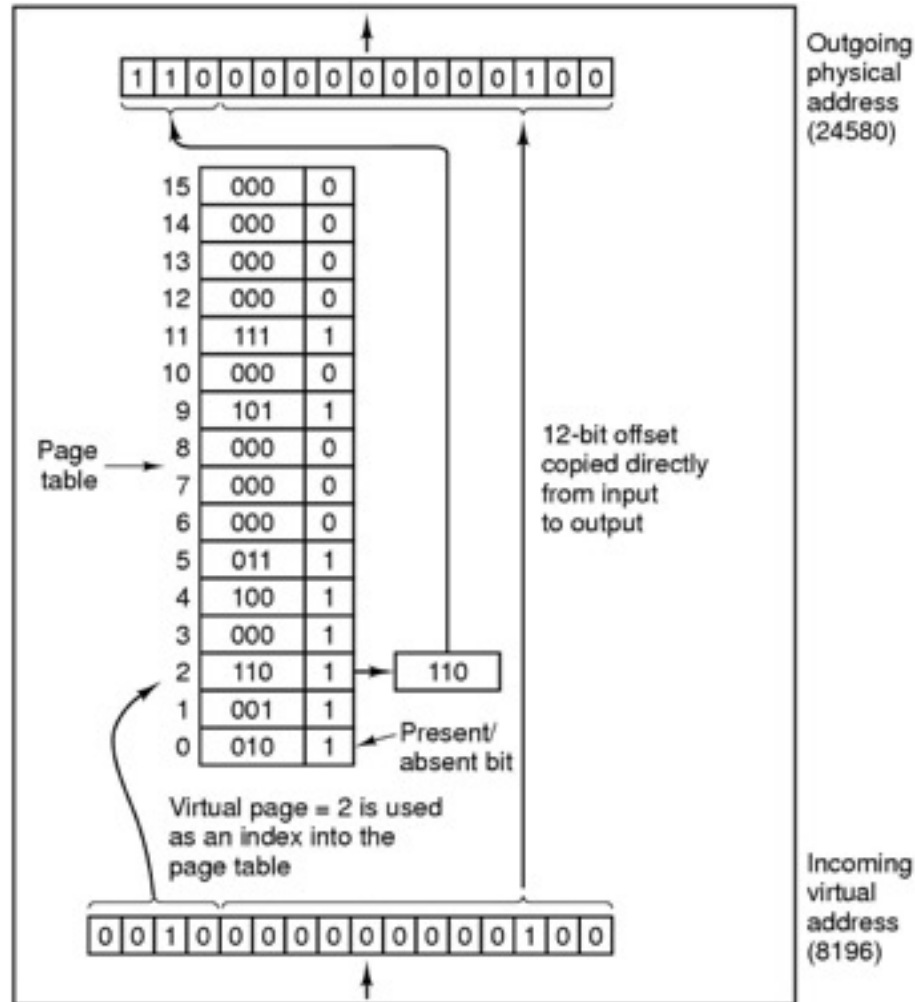
- ◆ MMU = Memory Management Unit
- ◆ Part of Hardware that accompanies the CPU
- ◆ Converts Virtual Addresses to Physical Addresses

Paging

The relation between virtual addresses and physical memory addresses given by page table

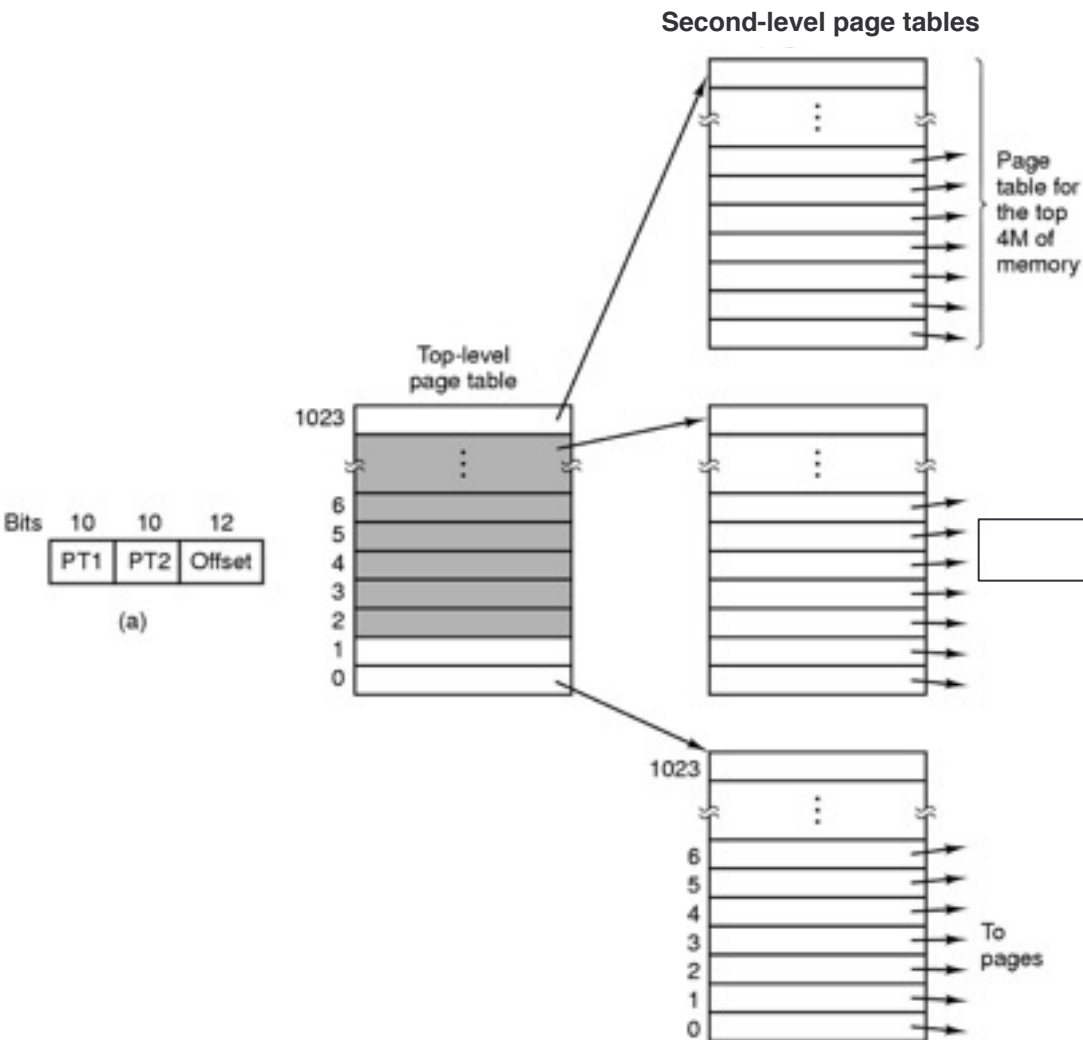


Page Tables (1)



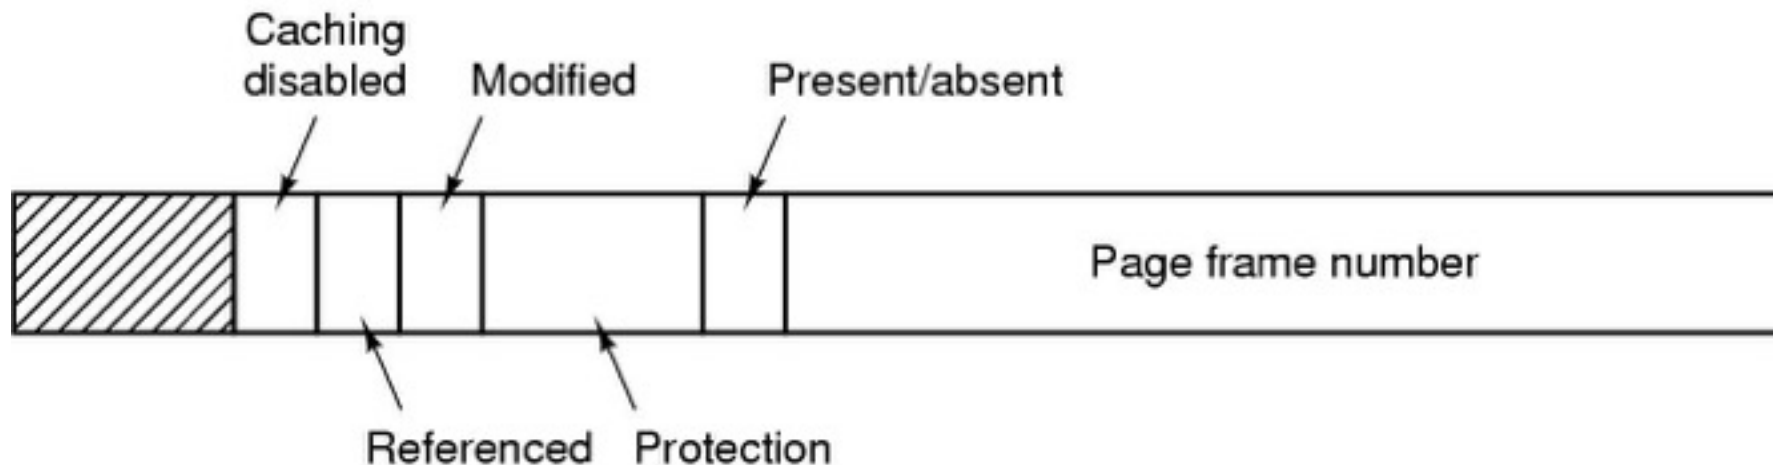
Internal operation of MMU with 16 4 KB pages

Page Tables (2)



- ◆ 32 bit address with 2 page table fields
- ◆ Two-level page tables
- ◆ PT too Big for MMU
 - Place it in main memory
- ◆ But how does MMU know where to find PT?
 - Registers (CR2 on Intel)

Typical page table entry



- ◆ Page Frame number = physical page number for the virtual page represented by the PTE
- ◆ Referenced bit: Whether the page was accessed since last time the bit was reset.
- ◆ Modified bit: Also called “Dirty” bit. Whether the page was written to, since the last time the bit was reset.
- ◆ Protection bits: Whether the page is readable? writeable? executable? contains higher privilege code/data?
- ◆ Present/Absent bit: Whether the PTE contains a valid page frame #. Used for marking swapped/unallocated pages.

TLBs – Translation Lookaside Buffers

Valid	Virtual page	Modified	Protection	Page frame
1	140	1	RW	31
1	20	0	R X	38
1	130	1	RW	29
1	129	1	RW	62
1	19	0	R X	50
1	21	0	R X	45
1	860	1	RW	14
1	861	1	RW	75

- ◆ TLB is a small cache that speeds up the translation of virtual addresses to physical addresses.
- ◆ TLB is part of the MMU hardware (comes with CPU)
- ◆ It is not a Data Cache or Instruction Cache. Those are separate.
- ◆ TLB simply caches translations from virtual page number to physical page number so that the MMU don't have to access page-table in memory too often.
- ◆ On x86 architecture, TLB has to be “flushed” upon every context switch because there is no field in TLB to identify the process context.

Practical, transparent operating system support for superpages

Juan Navarro • Sitaram Iyer
Peter Druschel • Alan Cox

Rice University

Overview

- ◆ Increasing cost in TLB miss overhead
 - growing working sets
 - TLB size does not grow at same pace
- ◆ Processors now provide superpages
 - one TLB entry can map a large region
- ◆ OSs have been slow to harness them
 - no transparent superpage support for apps
- ◆ This talk: a practical and transparent solution to support superpages

Translation look-aside buffer

- ◆ TLB caches virtual-to-physical address translations
- ◆ TLB coverage
 - amount of memory mapped by TLB
 - amount of memory that can be accessed without TLB misses

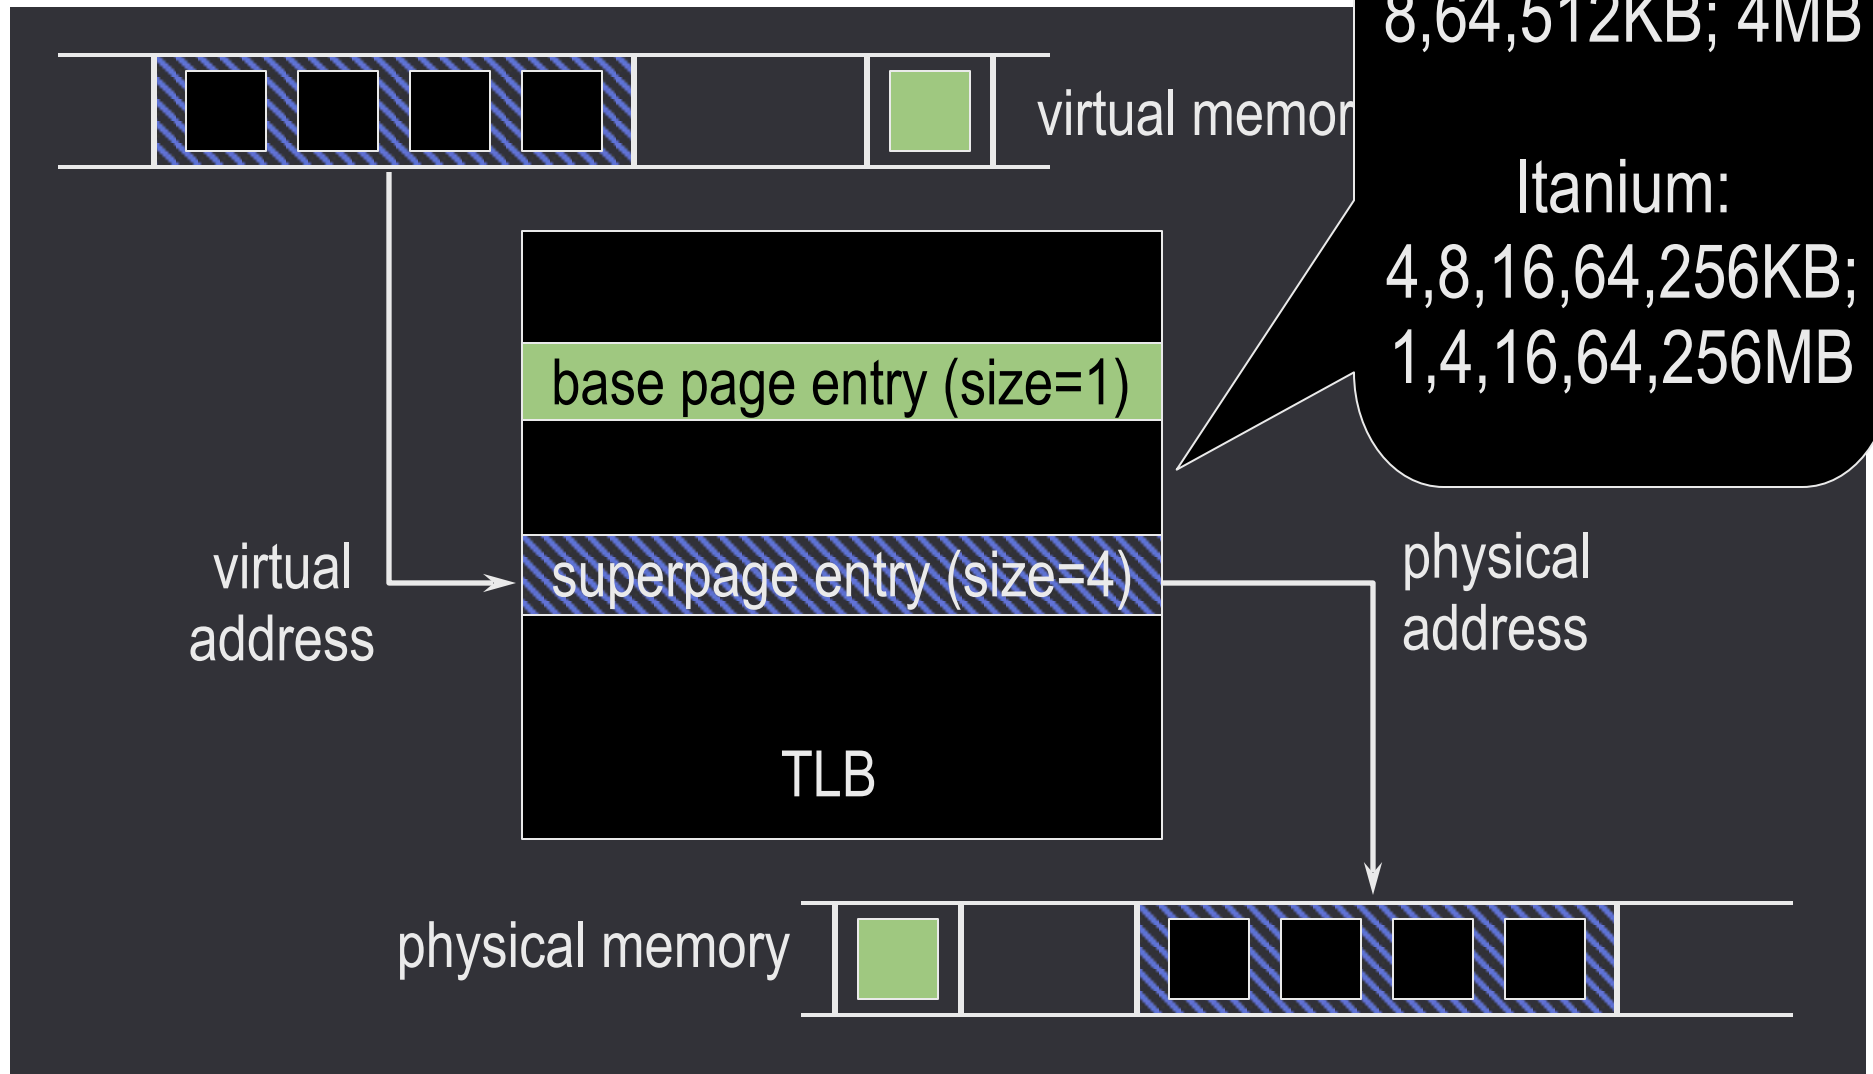
How to increase TLB coverage

- ◆ Typical TLB coverage \approx 1 MB
- ◆ Use superpages!
 - large and small pages
 - Increase TLB coverage
 - no increase in TLB size

What are these superpages anyway?

- ◆ Memory pages of larger sizes
 - supported by most modern CPUs
- ◆ Otherwise, same as normal pages
 - power of 2 size
 - use only one TLB entry
 - contiguous
 - aligned (physically and virtually)
 - uniform protection attributes
 - one reference bit, one dirty bit

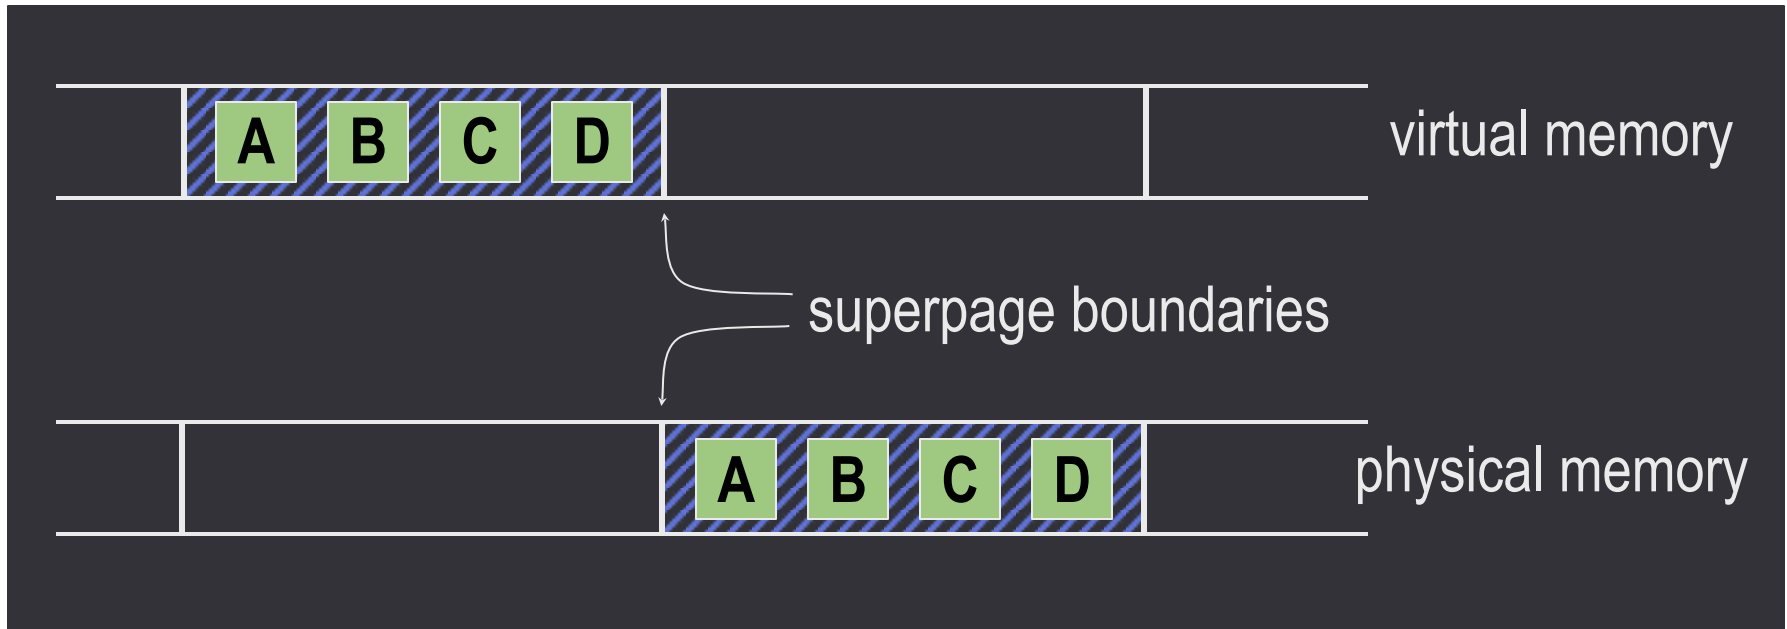
A superpage TLB



II

The superpage problem


Issue 1: superpage allocation



◆ How / when / what size to allocate?

Issue 2: promotion

- ◆ Promotion: create a superpage out of a set of smaller pages
 - mark page table entry of each base page
- ◆ When to promote?



Create sm
May inc

Forcibly populate pages?
May incur I/O cost or increase
internal fragmentation.

to touch pages?
opportunity to increase
coverage.

Issue 3: demotion

Demotion: convert a superpage into smaller pages

- ◆ when page attributes of base pages of a superpage become non-uniform
- ◆ during partial pageouts

Issue 4: fragmentation

- ◆ Memory becomes fragmented due to
 - use of multiple page sizes
 - scattered *wired* (non-pageable) pages
- ◆ Contiguity: contended resource
- ◆ OS must
 - use contiguity restoration techniques
 - trade off impact of contiguity restoration against superpage benefits

Previous approaches

- ◆ Reservations
 - one superpage size only
- ◆ Relocation
 - move pages at promotion time
 - must recover copying costs
- ◆ Eager superpage creation (IRIX, HP-UX)
 - size specified by user: non-transparent
- ◆ Hardware support
 - Contiguous virtual superpage mapped to discontiguous physical base pages
- ◆ Demotion issues not addressed
 - large pages partially dirty/referenced

III Design

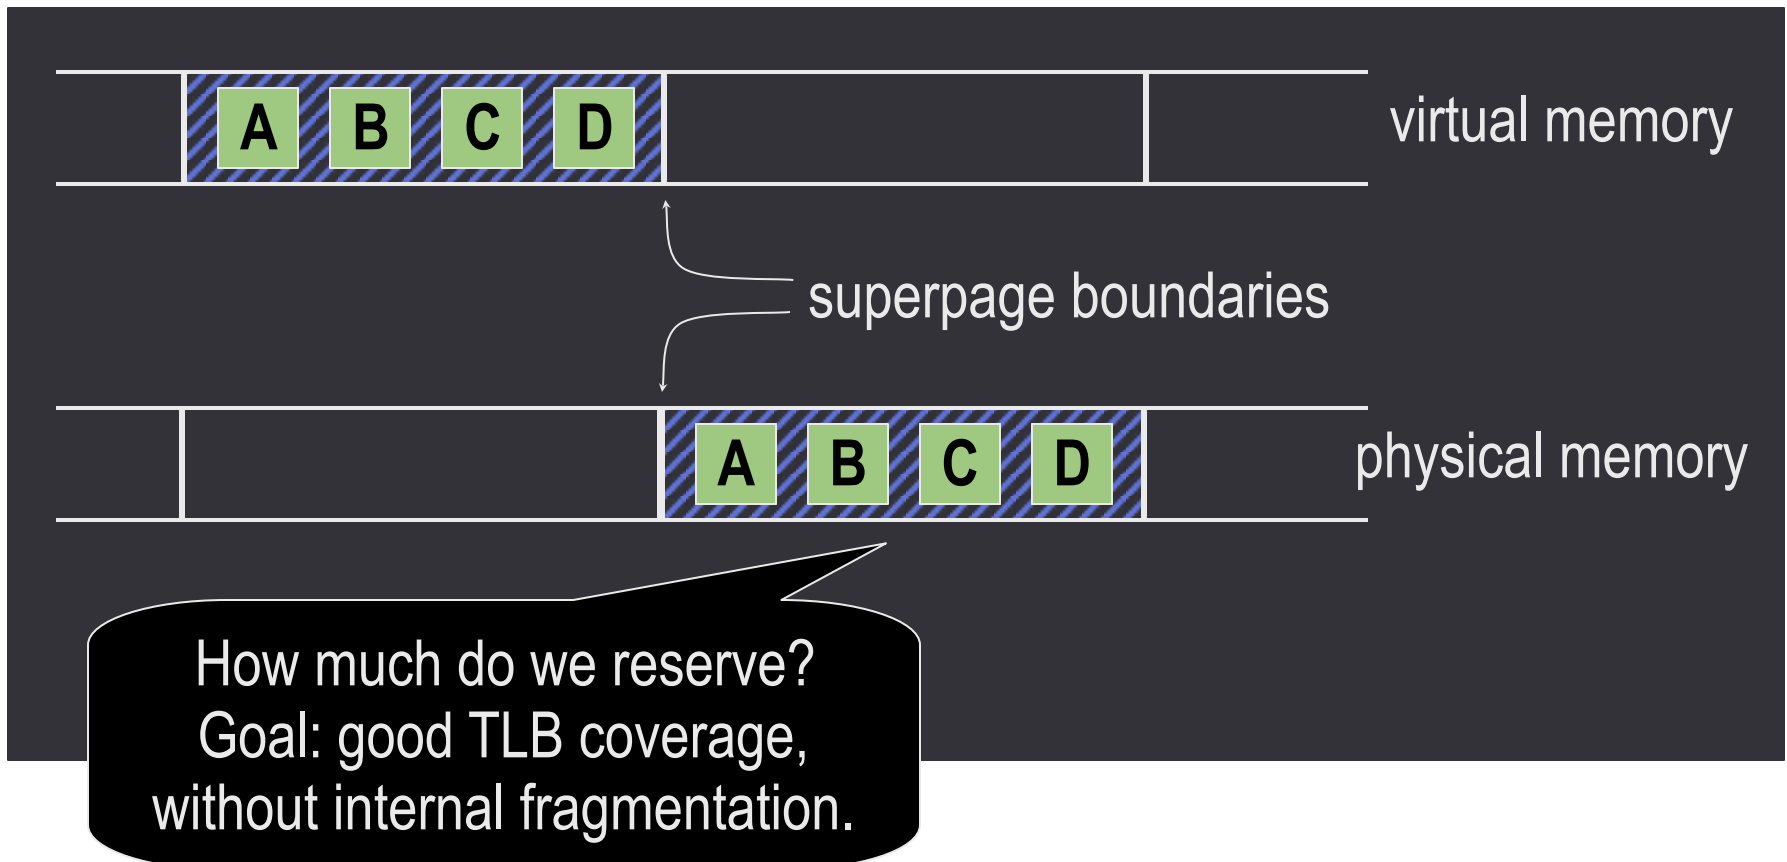
Key observation

Once an application touches the first page of a memory object then it is likely that it will quickly touch every page of that object

- ◆ Example: array initialization
- ◆ Opportunistic policies
 - superpages as large and as soon as possible
 - as long as no penalty if wrong decision

Superpage allocation

Preemptible reservations

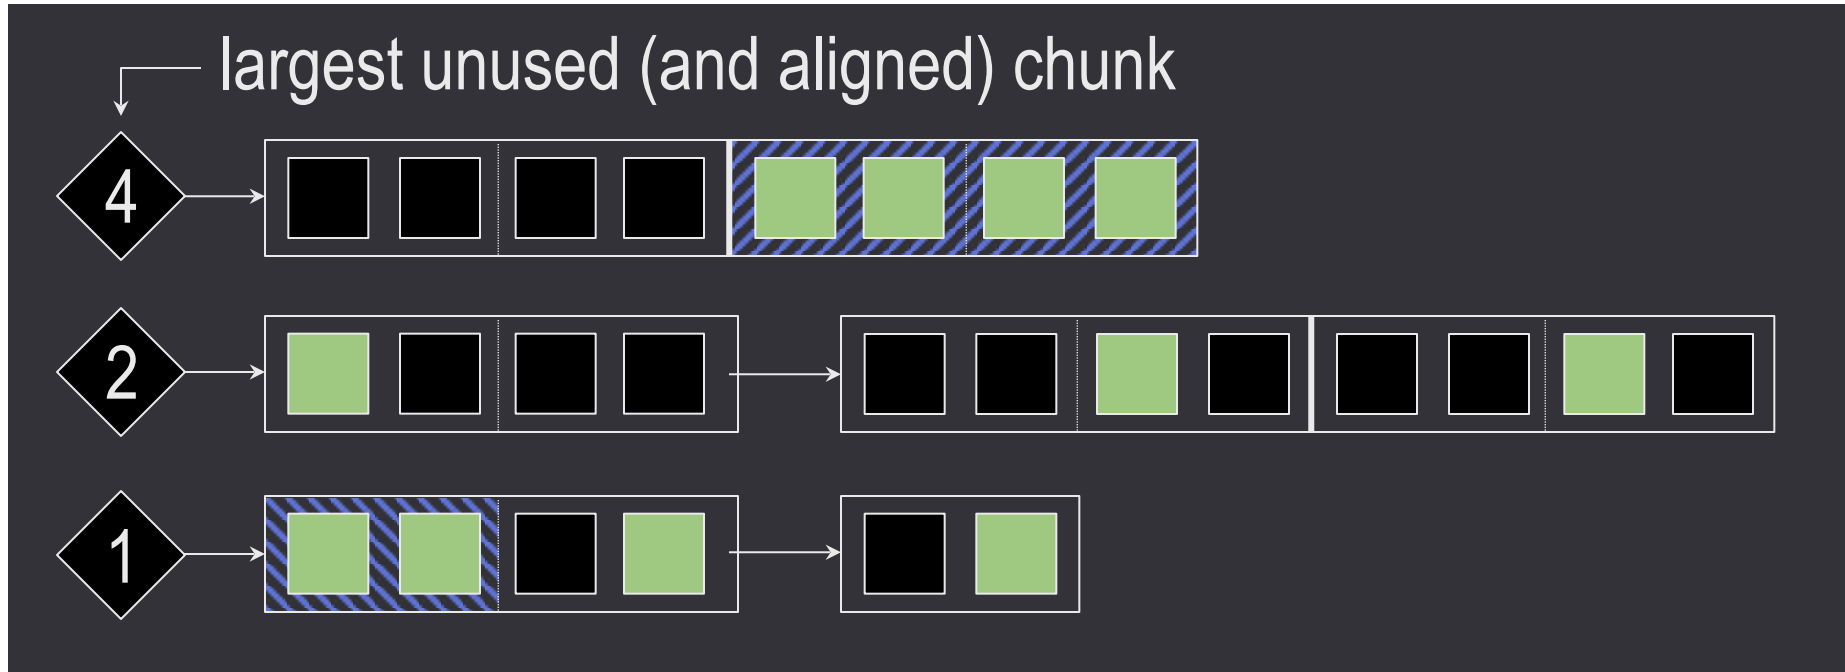


Allocation: reservation size

Opportunistic policy

- ◆ Go for biggest size that is no larger than the memory object (e.g., file)
- ◆ If required size not available, try preemption before resigning to a smaller size
 - preempted reservation had its chance

Allocation: managing reservations



best candidate for preemption at front:

- ◆ reservation whose most recently populated frame was populated the least recently

Incremental promotions

Promotion policy: opportunistic




Speculative demotions

- ◆ One reference bit per superpage
 - How do we detect portions of a superpage not referenced anymore?
- ◆ On memory pressure, demote superpages when resetting ref bit
- ◆ Re-promote (incrementally) as pages are referenced
- ◆ Demote also when the page daemon selects a base page as a victim page.

Demotions: dirty superpages

- ◆ One dirty bit per superpage
 - what's dirty and what's not?
 - page out entire superpage
- ◆ Demote on first write to clean superpage



- ◆ Re-promote (incrementally) as other pages are dirtied

Fragmentation control

- ◆ Low contiguity: modified page daemon for victim selection
 - restore contiguity
 - move clean, inactive pages to the free list
 - minimize impact
 - prefer pages that contribute the most to contiguity

- ◆ Cluster wired pages

IV

Experimental evaluation

Experimental setup

- ◆ FreeBSD 4.3
- ◆ Alpha 21264, 500 MHz, 512 MB RAM
- ◆ 8 KB, 64 KB, 512 KB, 4 MB pages
- ◆ 128-entry DTLB, 128-entry ITLB
- ◆ Unmodified applications

Best-case benefits

- ◆ TLB miss reduction usually above 95%
- ◆ SPEC CPU2000 integer
 - 11.2% improvement (0 to 38%)
- ◆ SPEC CPU2000 floating point
 - 11.0% improvement (-1.5% to 83%)
- ◆ Other benchmarks
 - FFT (200^3 matrix): 55%
 - 1000x1000 matrix transpose: 655%
- ◆ 30%+ in 8 out of 35 benchmarks

Why multiple superpage sizes

	64KB	512KB	4MB	All
FFT	1%	0%	55%	55%
galgel	28%	28%	1%	29%
mcf	24%	31%	22%	68%

Improvements with only one superpage size vs. all sizes

Conclusions

- ◆ Superpages: 30%+ improvement
 - transparently realized; low overhead
- ◆ Contiguity restoration is necessary
 - sustains benefits; low impact
- ◆ Multiple page sizes are important
 - scales to very large superpages

More info at

www.cs.rice.edu/~jnavarro/superpages