

Threads

Operating Systems

Kartik Gopalan


Chapter 2 Modern Operating Systems, Andrew Tanenbaum

Chapters 26 and 27, OSTEP book

Chapter 11 Advanced Programming in Unix Environment, By
Richard Stevens


If you want to do one task

- Start one process




If you want to do two task “concurrently”

- Start two processes
 - Maybe P1 forks P2
 - and P3...PN etc if more than two tasks
- Problem:
 - fork is expensive
 - cold-start penalty



If P1 and P2 want to talk to each other?

- E.g. access the same data or synchronize?
- Two different address spaces
 - Need to use IPC
 - shared memory, pipes, sockets, signals
- Problem
 - kernel transitions are expensive
 - May need to copy data
 - user—>kernel—>user
 - Inter-process Shared memory is a pain to set up.



Option 1:Event-driven programming


- Make one process do all the tasks
- Busy loop polls for events and executes tasks for each event
- No IPC needed
- Length of the busy loop determines response latency
- Stateful event responses complicate the code
 - What if i^{th} occurrence of event 1 effects the j^{th} event processing ?

P1

```
while(1)
{
    if (event 1) do task 1;
    if (event 2) do task 2;
    ...
    if (event N) do task N;
}
```

Option 2: Use threads


- Multiple threads of execution per process
- Each thread has its own
 - Program counter
 - Stack, stack pointer
 - Registers
- All threads share
 - one virtual address space
 - code, heap and static data
- Lower context switching overhead
- No IPC
 - Zero data transfer cost
 - Only need inter-thread synchronization



Other Shared and non-shared components

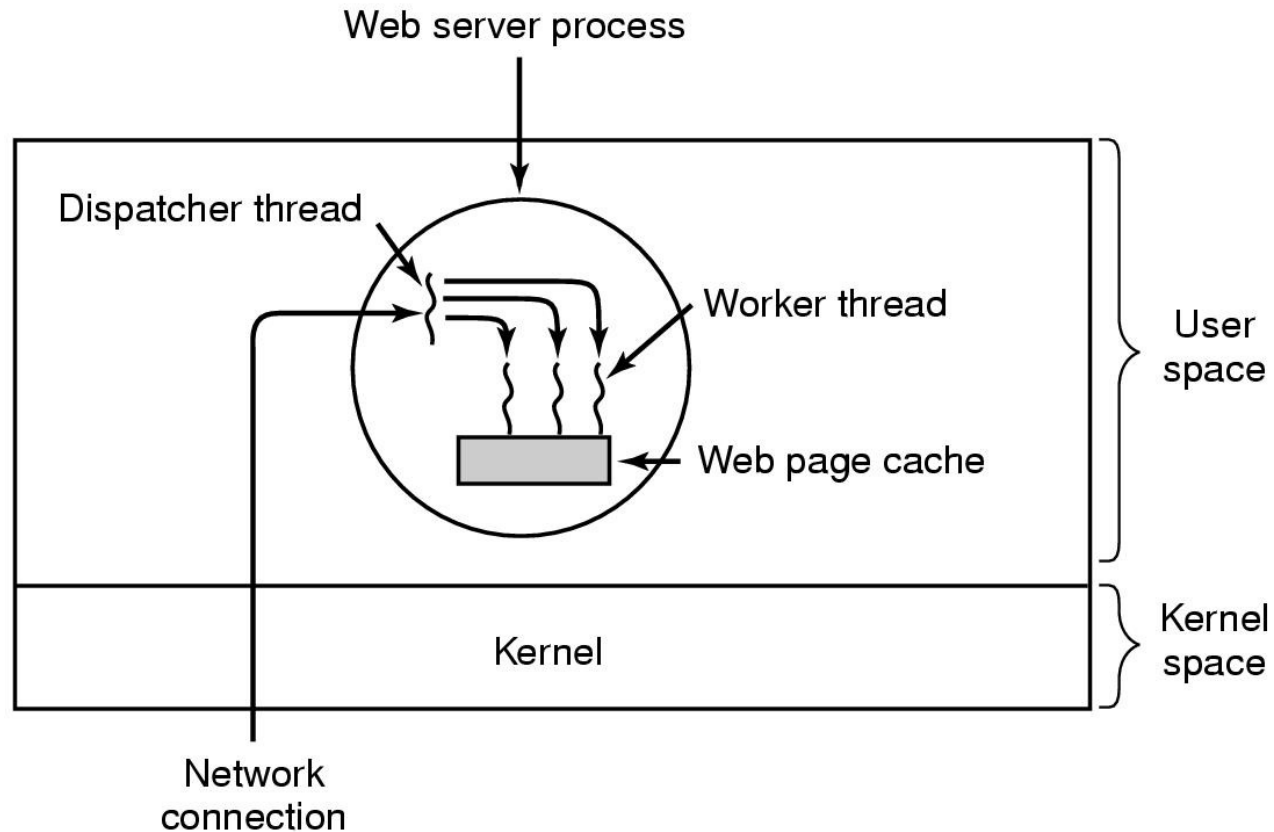
- Shared components
 - Open descriptors (files, sockets etc)
 - Signals and Signal handlers
- Not shared
 - Thread ID
 - Errno
 - Priority

Example: A word processor with three threads



- First thread handles keyboard input
- Second thread handles screen display
- Third thread handles saving the document to disk

Example: a multi-threaded web server




- A dispatcher thread waits for and accepts network connections
- Several worker threads
 - Each worker processes one network connection concurrently

Disadvantages of Threads


- Shared State!
 - Global variables are shared between threads.
 - Accidental data changes can cause errors.
- Threads and signals don't mix well
 - Common signal handler for all threads in a process
 - Which thread to signal? Everybody!
 - Royal pain to program correctly.
- Lack of robustness
 - Crash in one thread will crash the entire process.
- Some library functions may not be thread-safe
 - Library Functions that return pointers to static internal memory. E.g. `gethostbyname()`
 - Less of a problem these days.

Two types of threads: user-level and kernel-level



User-level threads


- User-level libraries provide multiple threads,
- OS kernel does not recognize user-level threads
- Threads execute when the process is scheduled



Kernel-level threads

- OS kernel provides multiple threads per process
- Each thread is scheduled independently by the kernel's CPU scheduler


Hybrid Implementations



Multiplexing user-level threads within each kernel-level threads

Local Thread Scheduling


- Next thread is picked from among the threads belonging to the current process
- Each process gets a timeslice from kernel.
- Then the timeslice is divided up among the threads within the current process
- Local scheduling can be implemented with either
 - Kernel-level threads OR
 - User-level threads.
- Scheduling decision requires only local knowledge of threads within the current process.



- For example, say process timeslice may be 50ms, and each thread within the process runs for 5 msec/CPU burst

Global Thread scheduling

- Next thread to be scheduled is picked up from ANY process in the system.
 - Not just the current process
- Timeslice is allocated at the granularity of threads
 - No notion of per-process timeslice
- Global scheduling can be implemented only with kernel-level threads
 - Picking the next thread requires global knowledge of threads in all processes.



Possible: A1, A2, A3, A1, A2, A3

Also possible: A1, B1, A2, B2, A3, B3

- For example each thread runs for 10msec per CPU burst

Thread Creation and termination

- Creation
 - `int pthread_create(pthread_t * thread, pthread_attr_t * attr, void * (*start_routine)(void *), void * arg);`
- Two ways to perform thread termination
 1. Return from initial function.
 2. `void pthread_exit(void * status)`
- Waiting for child thread in parent
 - `pthread_join(...)`
 - equivalent to `waitpid`

Threaded program - example

```
// shared counter to be incremented by each thread
int counter = 0;

main()
{
    pthread_t tid[N];

    for (i=0;i<N;i++) {

        /*Create a thread in thread_func routine*/
        Pthread_create(&tid[i], NULL, thread_func, NULL);
    }

    for(i=0;i<N;i++)
        /* wait for child thread */
        Pthread_join(tid[i], NULL);
}

void *thread_func(void *arg)
{
    /* unprotected code - race condition*/
    counter = counter + 1;

    return NULL; // thread dies upon return
}
```


pthread synchronization operations

- Mutex operation
 - **pthread_mutex_init(...)**
 - **pthread_mutex_lock(...)**
 - **pthread_mutex_unlock (...)**
 - **pthread_mutex_trylock (...)**
- Condition variables
 - pthread_cond_wait (...)
 - pthread_cond_signal (...)
 - pthread_cond_broadcast (...)
 - pthread_cond_timedwait (...)